




Amirkabir University of Technology
(Tehran Polytechnic)
Department of Computer Engineering


Operating Systems Course

Hamid R. Zarandi
h_zarandi@aut.ac.ir


Textbooks



Silberschatz et. al. *Operating System Concepts*
9th Edition, Wiley, 2018.



Stallings, *Operating Systems: Internal and Design Principles*
9th Edition, Pearson, 2018.



Tanenbaum, *Modern Operating Systems*
4th Edition, Prentice-Hall, 2014.

Course highlights & grading

➤ Grading details

- Mid-term+Quiz (6)
- Final + Quiz (9)
- Homeworks (3)
- Project(s) (1)
- Presence (1)


➤ TA

- Mahshid Shiri
- Matin Jafar-Aghaei
- Soroush Mortazavi-Moghadam

➤ Highlights

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Structure
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management
- Protection and Security
- Kernel Data Structures
- Computing Environments
- Open-Source Operating Systems

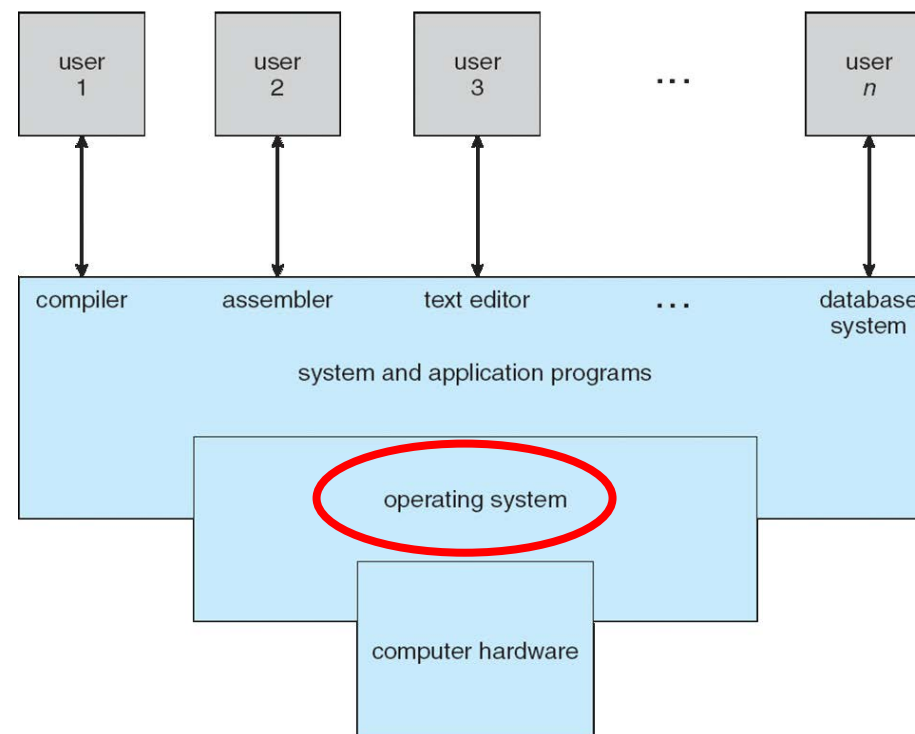
OS topics



Introductions to Operating Systems

What is an Operating System?

- A **program** that acts as an **intermediary** between a **user** of a computer and the **computer hardware**
 - User can execute programs **conveniently** & **efficiently**



What operating systems do?

- Users want convenience, **ease of use** and **good performance**
 - Don't care about **resource utilization**
- But shared computer such as **mainframe** or **minicomputer** must keep all users happy
- Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- Handheld computers are resource poor, optimized for **usability** and **battery life**
- Some computers have **little or no user interface**, such as embedded computers in devices and automobiles

Operating system definition

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use

- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer

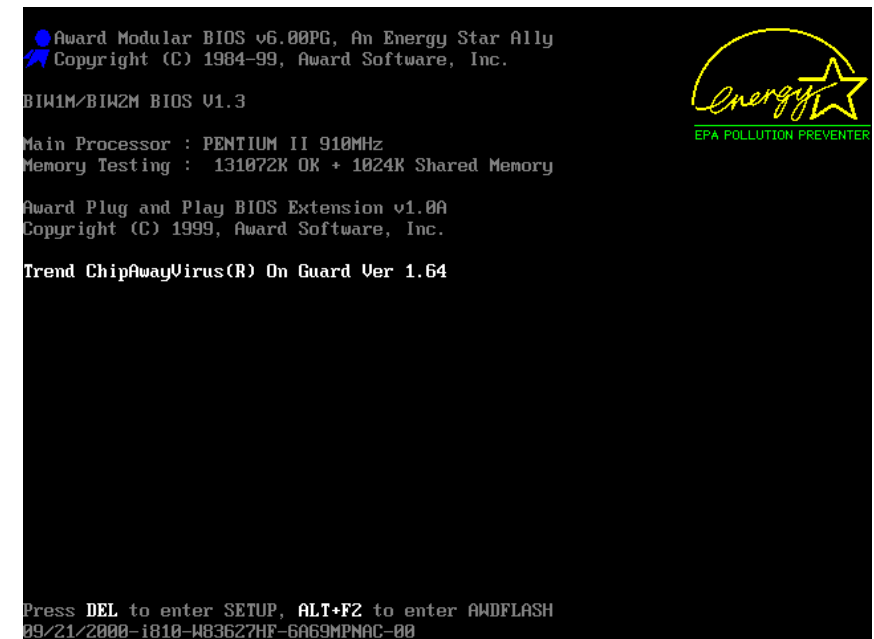


No uniform definition!

- No universally accepted definition!
- “The one program running at all times on the computer” is the kernel.

Computer startup


- **Bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, generally known as **firmware**
 - Initializes all aspects of system
 - Loads operating system kernel and starts execution




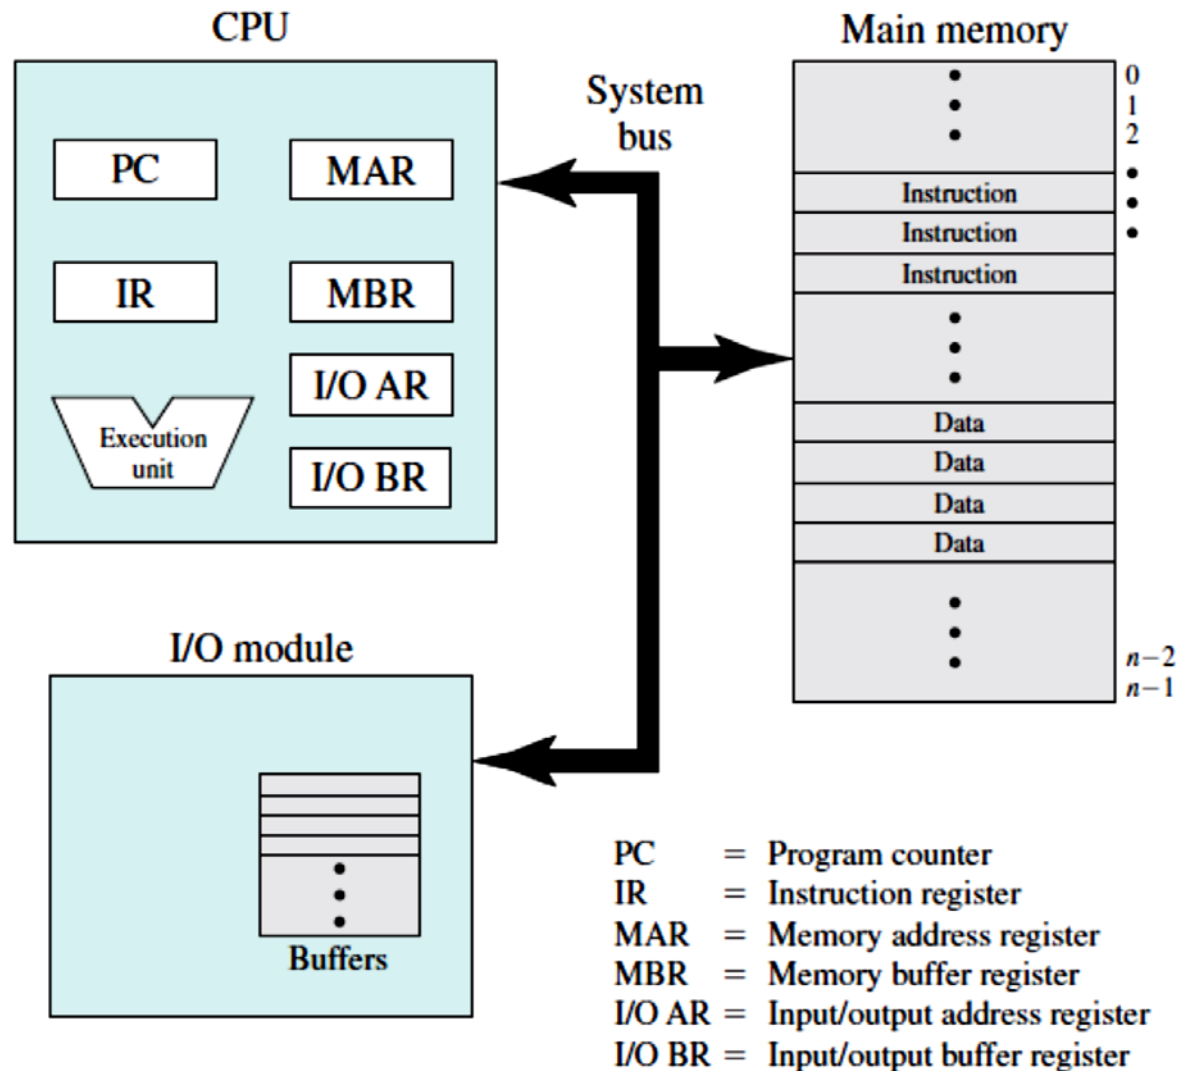
Computer system organization

➤ Computer-system operation

- One or more CPUs, device controllers connect through common bus providing access to shared memory
- Concurrent execution of CPUs and devices competing for memory cycles



Computer system, basic elements



Computer-system operation

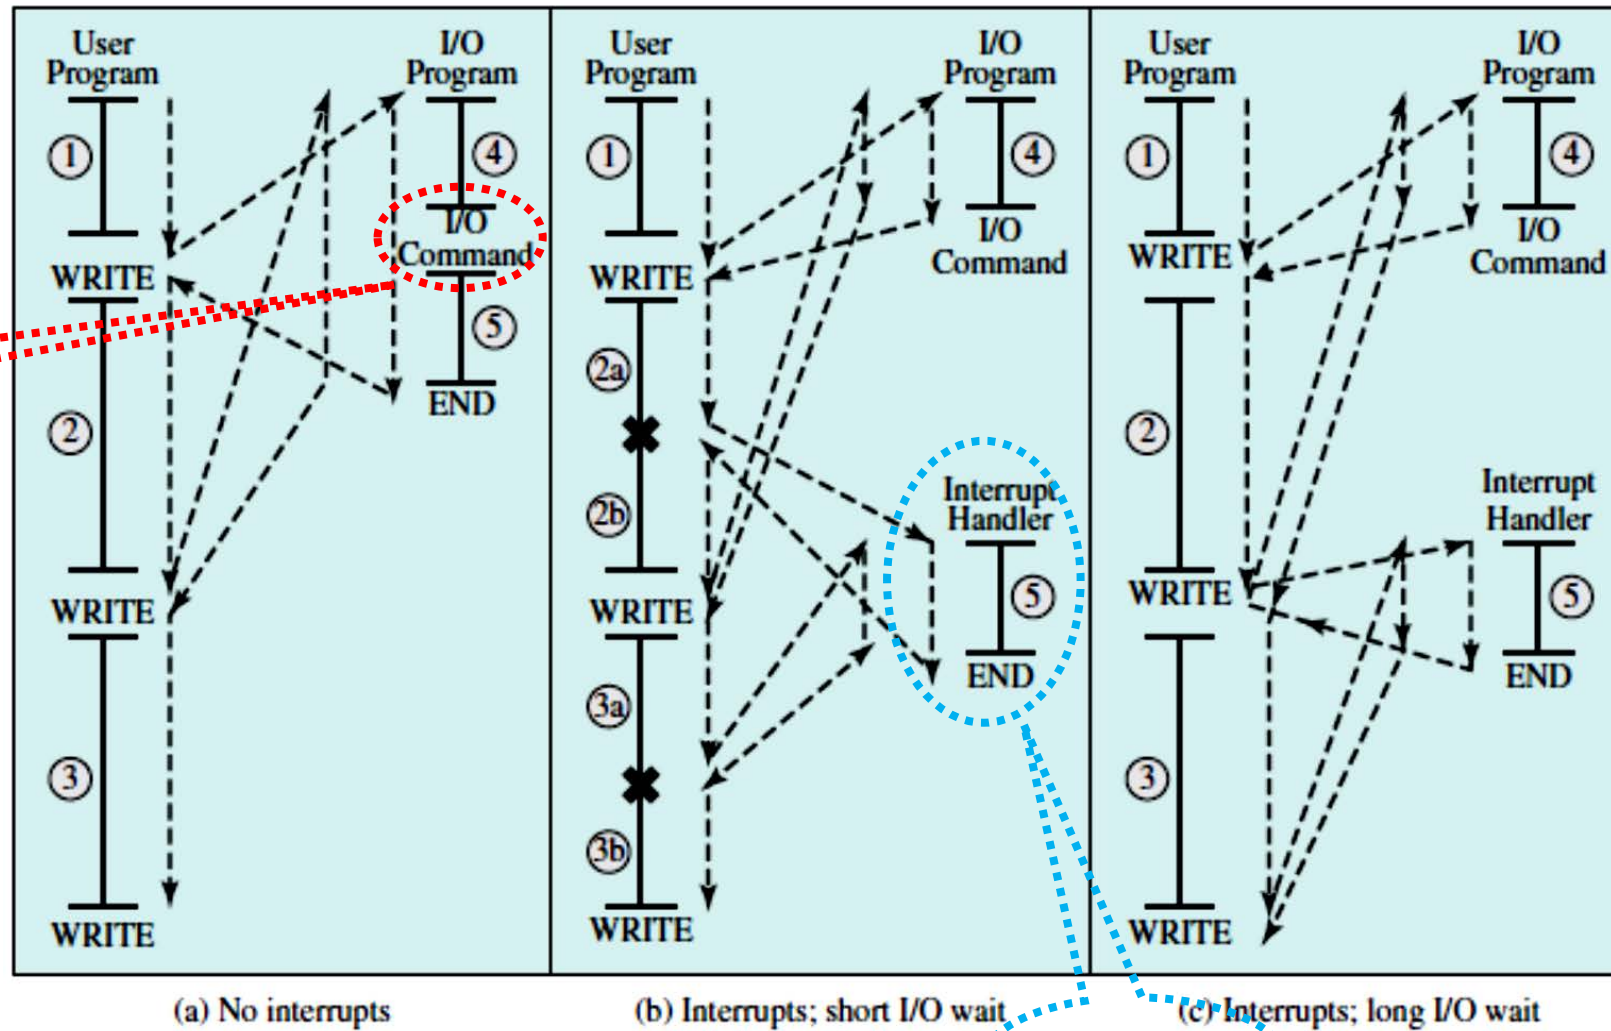
- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an **interrupt**

OS is interrupt driven!

Interrupts

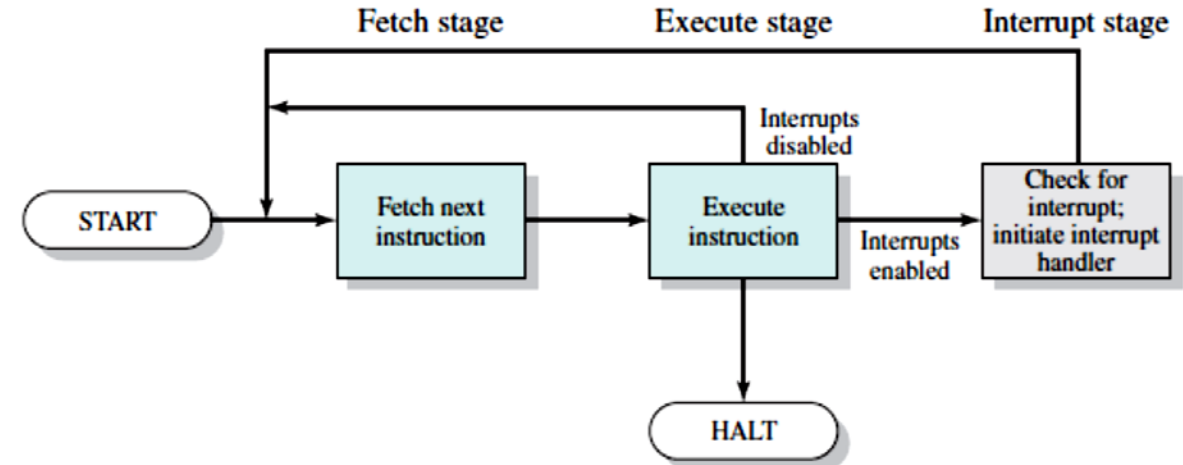
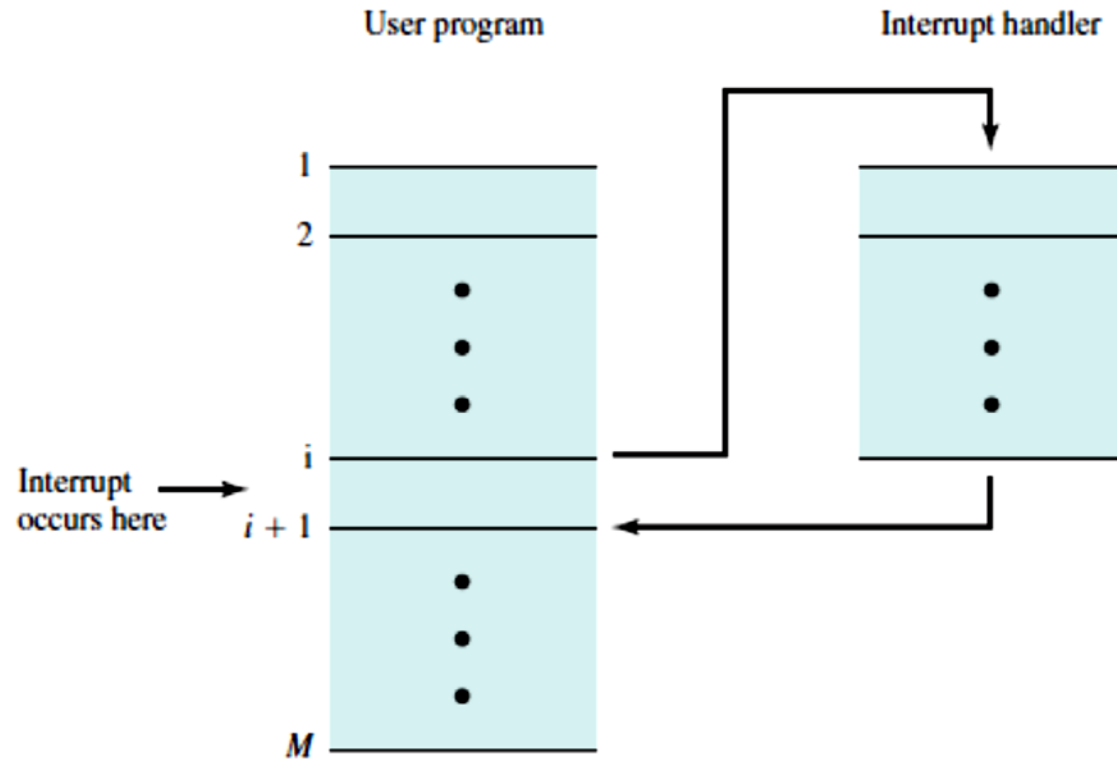
Program	Generated by some condition that occurs as a result of an instruction execution, such as arithmetic overflow, division by zero, attempt to execute an illegal machine instruction, and reference outside a user's allowed memory space.
Timer	Generated by a timer within the processor. This allows the operating system to perform certain functions on a regular basis.
I/O	Generated by an I/O controller, to signal normal completion of an operation or to signal a variety of error conditions.
Hardware failure	Generated by a failure, such as power failure or memory parity error.

Control flow (w/wo interrupts)

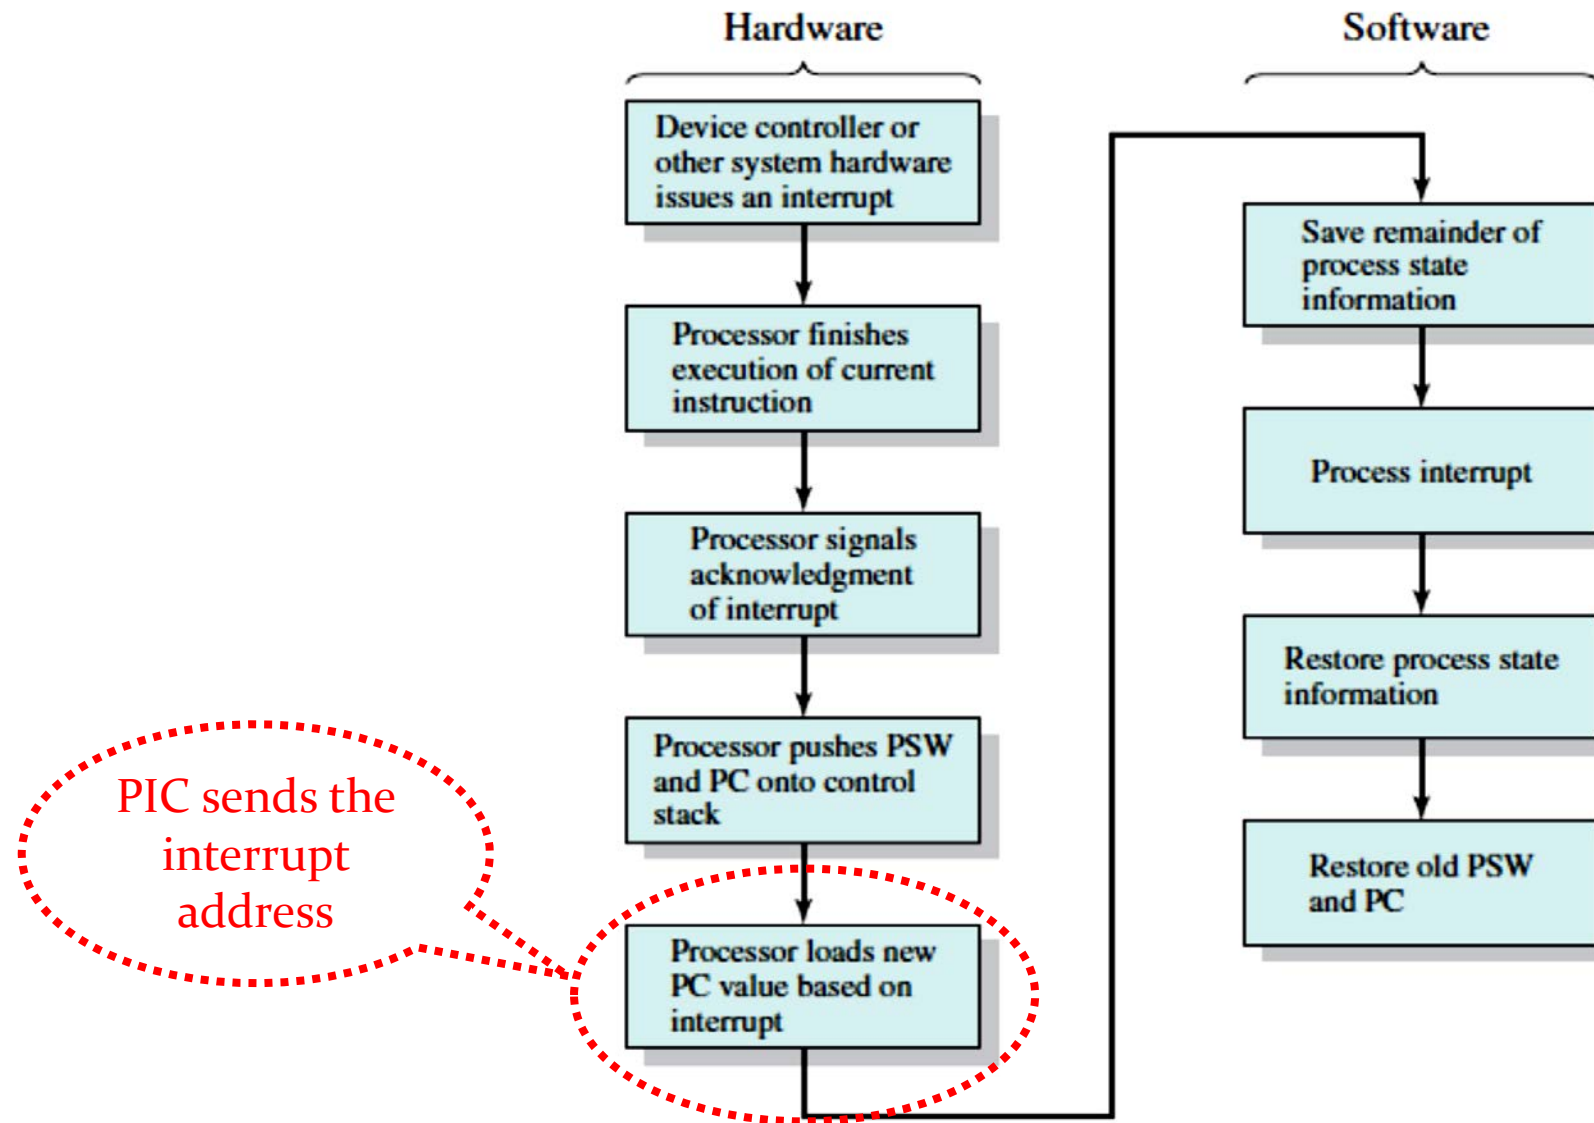


✕ = interrupt occurs during course of execution of user program


Handling interrupts




Interrupt handling



Multiple interrupts



(a) Sequential interrupt processing



(b) Nested interrupt processing

Disabling interrupt can be done to prevent future ones.


Common functions of interrupts

- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- **An operating system is interrupt driven**

Interrupt handling

- The operating system preserves the state of the CPU by storing registers and the program counter
- Determines which type of interrupt has occurred:
 - polling
 - vectored interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt

Interrupt timeline



Storage Definitions and Notation Review

Storage definitions and notation review

The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes.

A **kilobyte**, or **KB**, is 1,024 bytes

a **megabyte**, or **MB**, is $1,024^2$ bytes

a **gigabyte**, or **GB**, is $1,024^3$ bytes

a **terabyte**, or **TB**, is $1,024^4$ bytes

a **petabyte**, or **PB**, is $1,024^5$ bytes

Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).

Storage structure

- **Main memory** – only large storage media that the CPU can access directly
 - Random access
 - Typically **volatile**
- **Secondary storage** – extension of main memory that provides large **nonvolatile** storage capacity
- **Hard disks** – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer
- **Solid-state disks** – faster than hard disks, nonvolatile
 - Various technologies
 - Becoming more popular



Storage hierarchy

- **Storage systems organized in hierarchy**
 - Speed
 - Cost
 - Volatility

- **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage

- **Device Driver** for each device controller to manage I/O
 - Provides **uniform interface** between controller and kernel

Storage-device hierarchy




Performance of various levels of storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Movement between levels of storage hierarchy can be **explicit** or **implicit**


Migration of data “A” from disk to register

- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy

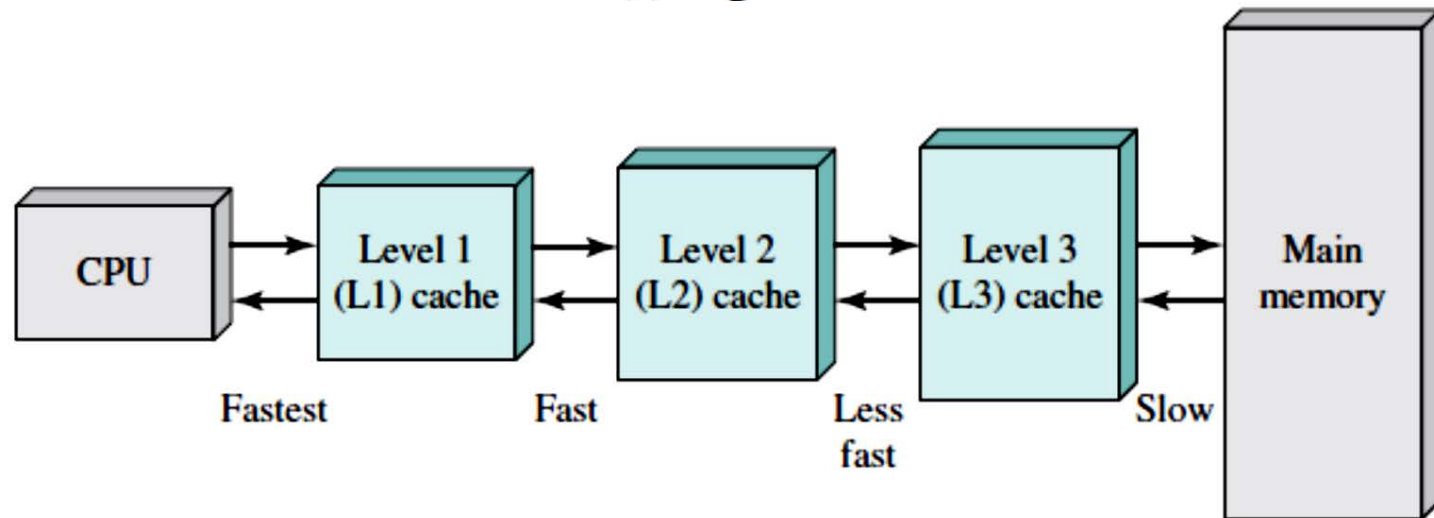


- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex
 - Several copies of a datum can exist
 - Various solutions covered in Chapter 17


Cache memory



(a) Single cache



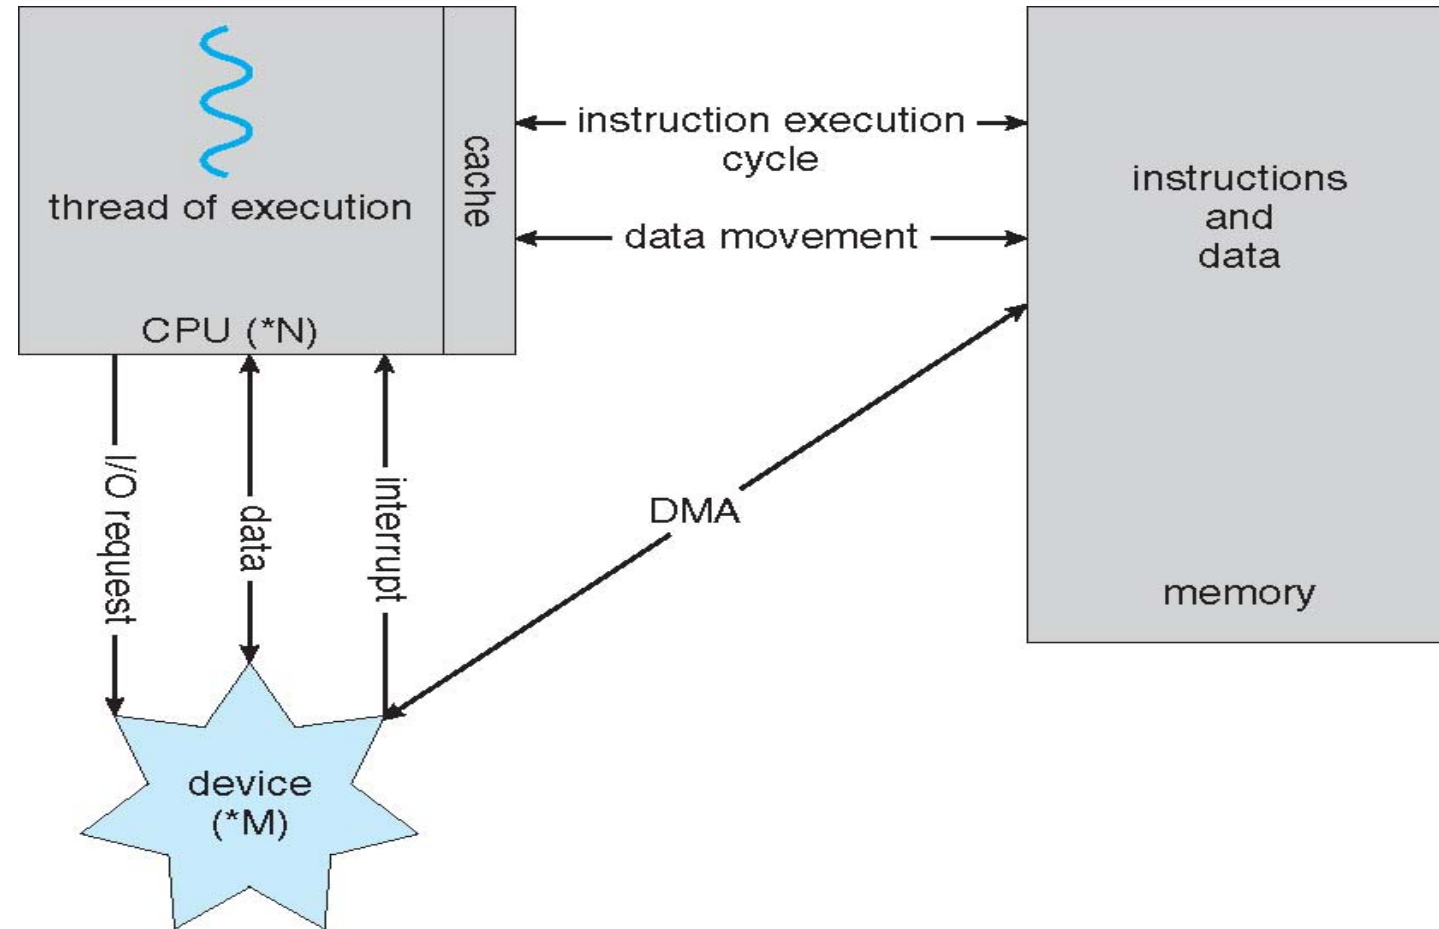
(b) Three-level cache organization



DMA: direct memory access

➤ IO operations

- Programmed IO (polling)
- Interrupt-driven
- DMA
 - **Whether** a read or write is requested
 - The **address of the I/O device**
 - The starting **location in memory** to read data/write data
 - The **number** of words




Computer System Architecture

Computer-system architecture

- **Most** systems use **a single general-purpose processor**
 - Most systems **have special-purpose processors** as well

- **Multiprocessors** systems growing in use and importance
 - Also known as **parallel systems, tightly-coupled systems**
 - Advantages include:
 1. **Increased throughput**
 2. **Economy of scale** – **multiprocessors** vs. **multiple single processor**
 3. **Increased reliability** – **graceful degradation** or **fault tolerance**
 - Two types:
 1. **Asymmetric Multiprocessing** – each processor is assigned a specific task (boss-worker).
 2. **Symmetric Multiprocessing** – each processor performs all tasks

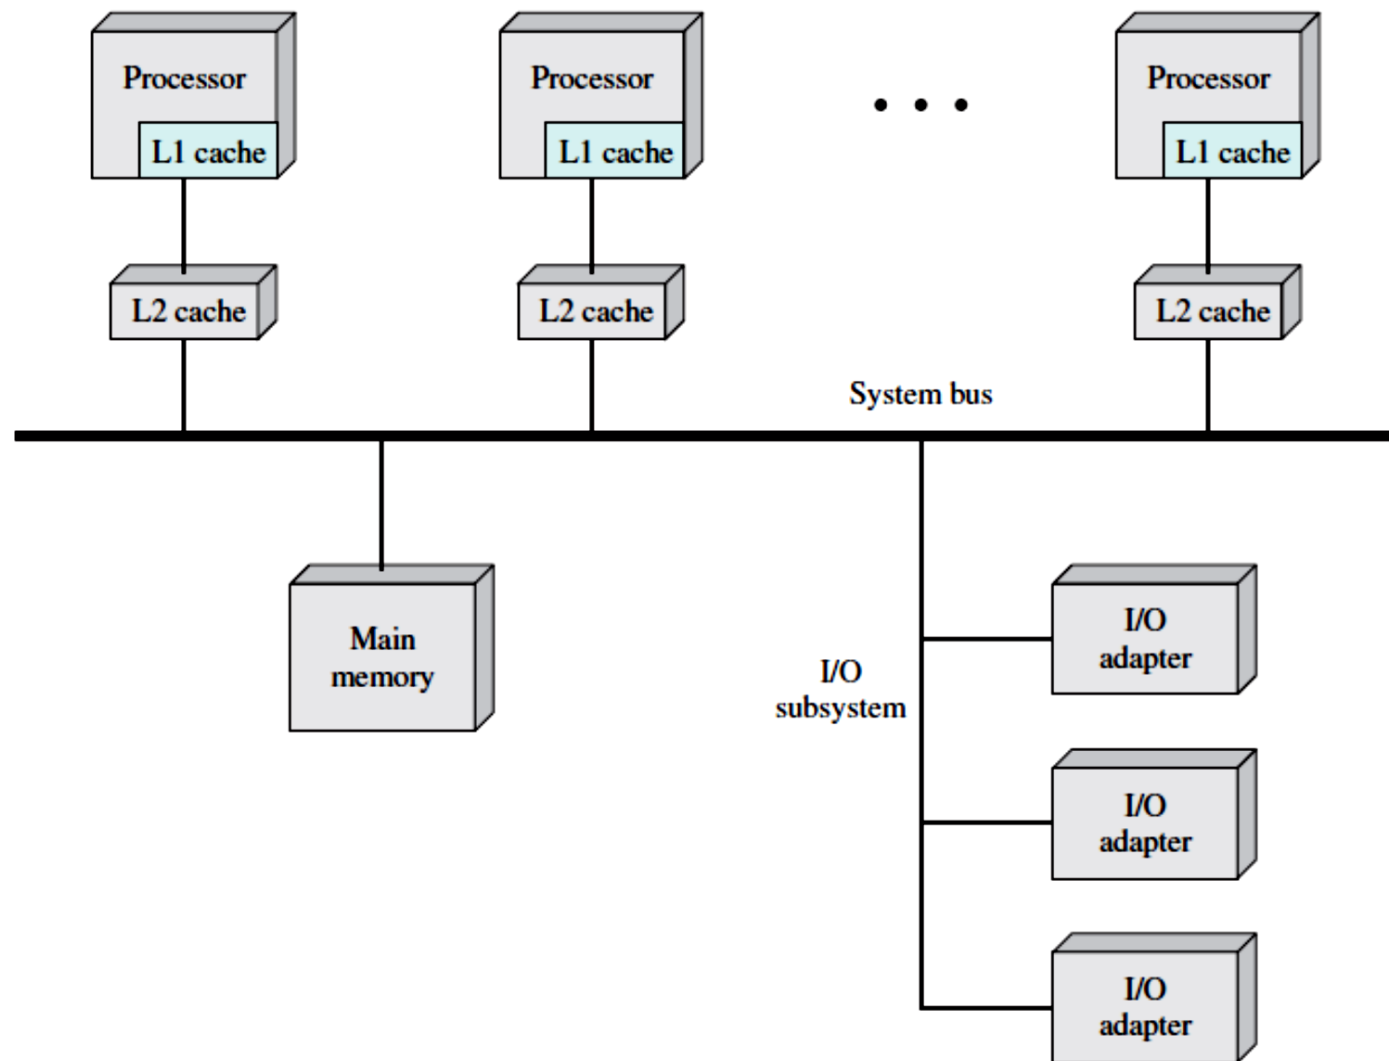
Symmetric multiprocessing architecture



➤ Multiprocessors memory access divisions:


- **UMA** (uniform memory access)
- **NUMA** (non-uniform memory access)

Symmetric multiprocessing organization

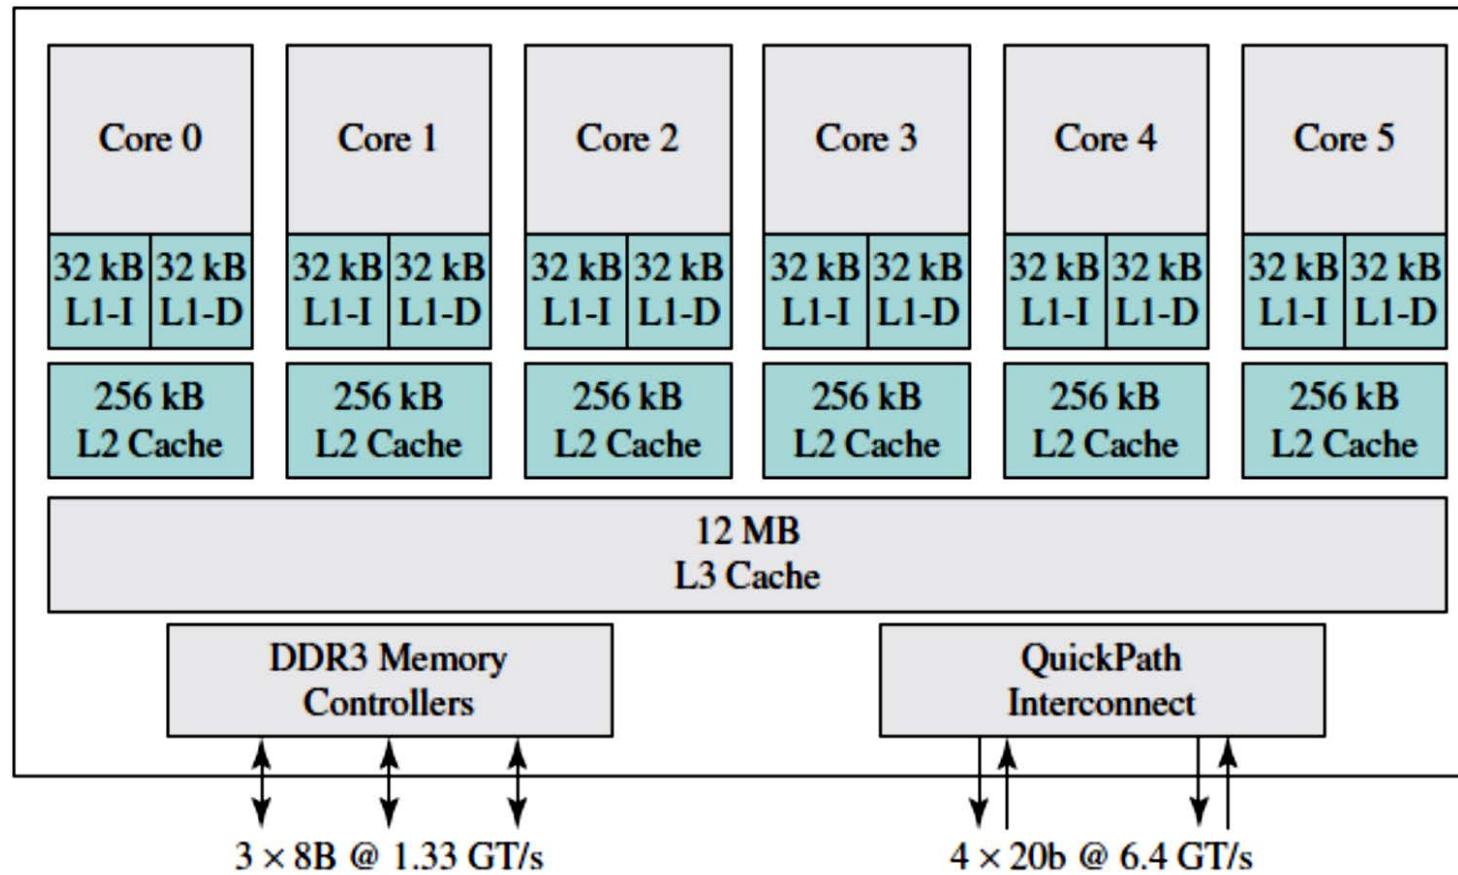


A dual-core design

- Multi-chip and **multicore**
- Systems containing all chips
 - Chassis containing multiple separate systems
- **Advantage:**
 - Faster communications
 - Less power consumption
- **Disadvantage:**
 - Performance gap of CPU utilization by software



Intel core i7-990X block diagram




Blade servers

- **Multiple** processor boards, I/O boards, and networking boards are placed in the **same chassis**.
- Each blade-processor
 - boots **independently**
 - runs **its own OS!**



Clustered systems

- Like multiprocessor systems, but multiple systems working together (**loosely-coupled systems**)
- Usually sharing storage via a **storage-area network (SAN)**
- Provides a **high-availability** service which survives failures
 - **Asymmetric clustering** has one machine in hot-standby mode
 - **Symmetric clustering** has multiple nodes running applications, monitoring each other
- Some clusters are for **high-performance computing (HPC)**
 - Applications must be written to use **parallelization**
- Some have **distributed lock manager (DLM)** to avoid conflicting operations





Operating System Structure

Operating system structure

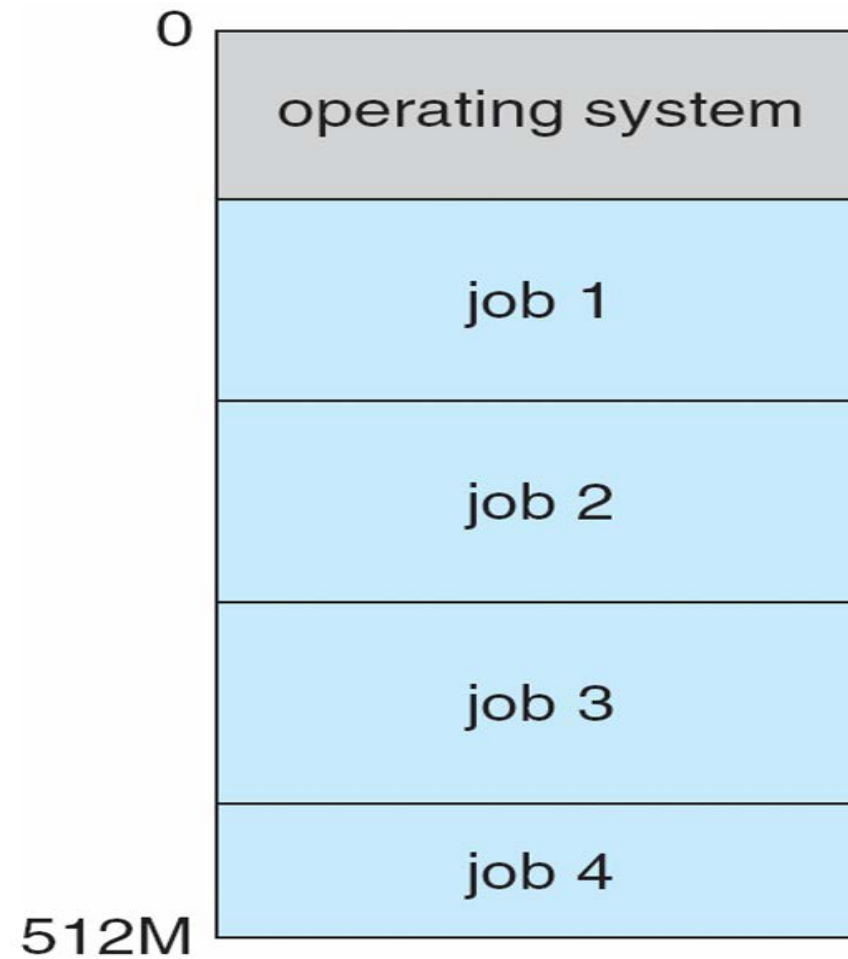
➤ Multiprogramming (Batch system) needed for efficiency

- Single user cannot keep CPU and I/O devices busy at all times
- Multiprogramming organizes jobs (code and data) so CPU always has one to execute
- A subset of total jobs in system is kept in memory
- One job selected and run via job scheduling
- When it has to wait (for I/O for example), OS switches to another job

➤ Timesharing (multitasking) is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating interactive computing

- Response time should be < 1 second
- Each user has at least one program executing in memory ⇒ process
- If several jobs ready to run at the same time ⇒ CPU scheduling
- If processes don't fit in memory, swapping moves them in and out to run
- Virtual memory allows execution of processes not completely in memory

Memory layout for multiprogrammed system




➤ **Job pool** on a disk!

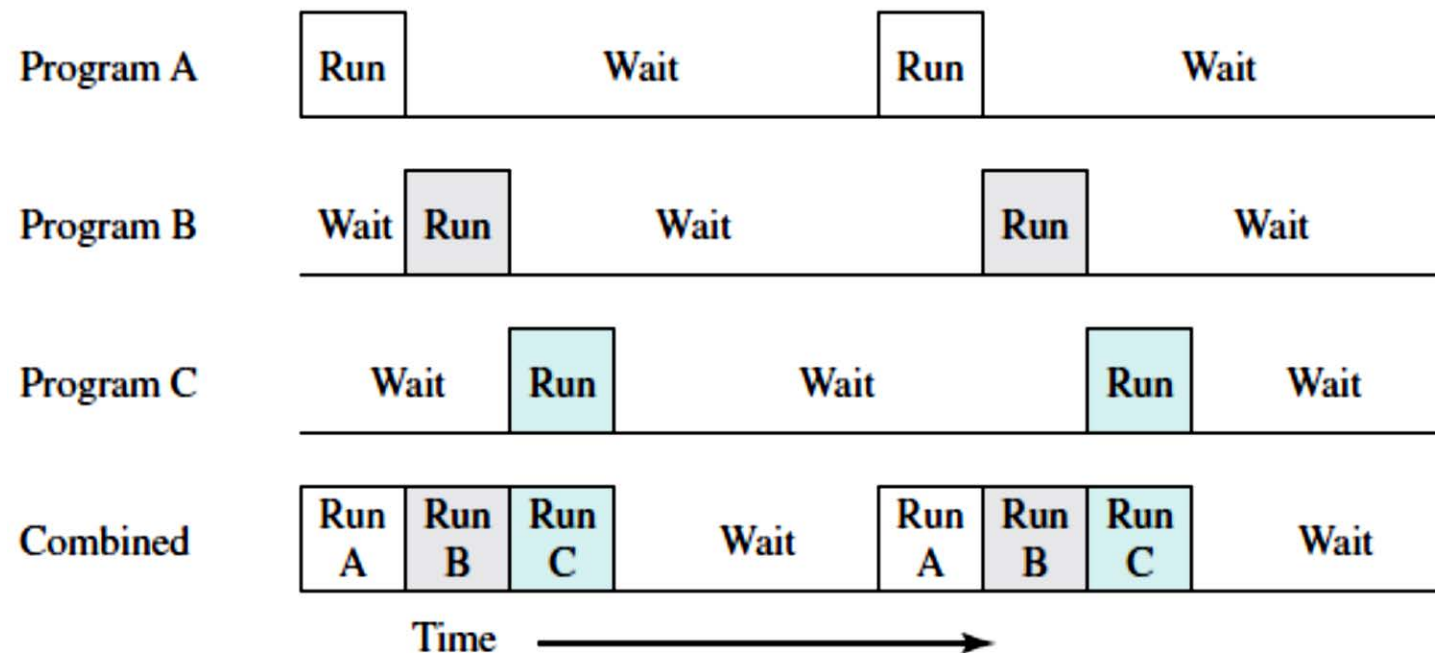
- Physical memory
- Virtual memory

System utilization example

Read one record from file	$15 \mu s$
Execute 100 instructions	$1 \mu s$
Write one record to file	$15 \mu s$
Total	$31 \mu s$
Percent CPU utilization = $\frac{1}{31} = 0.032 = 3.2\%$	



Increasing system utilization example



(c) Multiprogramming with three programs

➤ OS **idle state** is done by **HLT** assembly instruction (x86 opcode 0xF4)


Operating-system operations

- **Interrupt driven** (hardware and software)
 - **Hardware interrupt** by one of the devices
 - **Software interrupt** (**exception** or **trap**):
 - Software error (e.g., division by zero)
 - Request for operating system service
 - Other process problems include infinite loop, processes modifying each other or the operating system

Operating-system operations (cont.)

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user
- Increasingly CPUs support multi-mode operations
 - i.e. **virtual machine manager (VMM)** mode for guest **VMs**

Dual mode processors



Time management

- **Timer to prevent infinite loop / process hogging resources (Watchdog timer)**
 - Timer is set to interrupt the computer after some time period
 - Keep a counter that is decremented by the physical clock.
 - Operating system set the counter (privileged instruction)
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time

Process management

- A **process** is a program in *execution*. It is a unit of work within the system. **Program** is a *passive entity*, **process** is an *active entity*.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- **Multi-threaded process has one program counter per thread**
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes / threads

Process management activities

The operating system is responsible for the following activities in connection with process management:

- **Creating** and **deleting** both user and system processes
- **Suspending** and **resuming** processes
- Providing mechanisms for **process synchronization**
- Providing mechanisms for **process communication**
- Providing mechanisms for **deadlock handling**

Memory management

- To execute a program all (or part) of the instructions must be in memory
- All (or part) of the data that is needed by the program must be in memory.
- Memory management determines what is in memory and when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and de-allocating memory space as needed

Storage management

- OS provides uniform, logical view of information storage
 - **Abstracts** physical properties to logical storage unit - **file**
 - Each medium is controlled by device (i.e., disk drive, tape drive)
 - Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- File-System management
 - Files usually organized into **directories**
 - Access control on most systems to determine **who** can access **what**
 - OS activities include
 - **Creating** and **deleting** files and directories
 - **Primitives** to manipulate files and directories
 - **Mapping** files onto secondary storage
 - **Backup** files onto stable (non-volatile) storage media

Mass-storage management

- Usually disks used to store data that does **not fit in main memory** or data that must be kept for a “**long**” period of time
- Entire speed of computer operation hinges on **disk subsystem** and **its algorithms**
- OS activities
 - **Free-space management**
 - **Storage allocation**
 - **Disk scheduling**
- Some storage need not be fast
 - Tertiary storage includes **optical** storage, **magnetic** tape
 - Still must be managed – by **OS** or **applications**
 - Varies between **WORM** (write-once, read-many-times) and **RW** (read-write)

I/O subsystem

- One purpose of OS is to **hide peculiarities** of hardware devices from the user
- I/O subsystem responsible for:
 - **Memory management** of I/O including
 - **buffering** (storing data temporarily while it is being transferred)
 - **caching** (storing parts of data in faster storage for performance)
 - **spooling** (the overlapping of output of one job with input of other jobs)
 - General **device-driver** interface
 - **Drivers** for **specific** hardware devices



Protection and Security


Protection and security

- **Protection** – any mechanism for **controlling** access of **processes** or **users** to **resources** defined by the OS
- **Security** – defense of the system against **internal** and **external** attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine **who** can do **what**
 - User identities (**user IDs**, **security IDs**) include name and associated number, one per user
 - User ID then associated with all files, processes of that user to determine access control
 - Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
 - **Privilege escalation** allows user to change to effective ID with more rights


Kernel data structures

➤ Many similar to standard programming data structures


➤ *Singly linked list*



➤ *Doubly linked list*



➤ *Circular linked list*



Kernel data structures


➤ Binary search tree

left \leq right

- Search performance is $O(n)$
- **Balanced binary search tree** is $O(\lg n)$

➤ Example

- Linux CPU-scheduling algorithm




Kernel data structures

➤ Hash function can create a hash map

- Collision problem

Example: username, password



➤ Bitmap – string of n binary digits representing the status of n items (0011010111001010)

- Example: Disk block availability

➤ Linux data structures defined in

include files `<linux/list.h>`, `<linux/kfifo.h>`, `<linux/rbtree.h>`



Computing Environment

Computing environments (traditional)

- Stand-alone general purpose machines
- But blurred as most systems interconnect with others (i.e., the Internet)
- **Portals** provide web access to internal systems
- **Network computers** (**thin clients**) are like Web terminals
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous (**IoT**)– even home systems use **firewalls** to protect home computers from Internet attacks



Computing environments (Mobile)

- Handheld smartphones, tablets, etc
- What is the functional difference between them and a “traditional” laptop?
- Extra feature – more OS features (**GPS**, **gyroscope**)
- Allows new types of apps like *augmented reality*
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are **Apple iOS** and **Google Android**



Computing environments (distributed)

➤ Distributed computing


- Collection of separate, possibly heterogeneous, systems networked together

- **Network** is a communications path, **TCP/IP** most common

- ✓ Local Area Network (**LAN**)
- ✓ Wide Area Network (**WAN**)
- ✓ Metropolitan Area Network (**MAN**)
- ✓ Personal Area Network (**PAN**)

- **Network Operating System**


- Provides features between systems across network
- Communication scheme allows systems to exchange messages
- Illusion of a single system



Computing environments (Client-Server)


➤ Client-Server Computing

- ❑ Dumb terminals supplanted by smart PCs
- ❑ Many systems now **servers**, responding to requests generated by **clients**
 - ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
 - ▶ **File-server system** provides interface for clients to store and retrieve files



Computing environments (Peer-to-Peer)

- Another model of distributed system
- P2P does not distinguish **clients** and **servers**
 - **No bottleneck server!**
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - Registers its service with **central lookup service** on network, or
 - Broadcast request for service and respond to requests for service via **discovery protocol**
 - Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype



Computing environments (Virtualization)

- **Allows operating systems to run applications within other OSes**
 - Vast and growing industry


- **Emulation** used when source CPU type different from target type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code – **Interpretation**

- **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - **VMM** (virtual machine manager) provides virtualization services

Computing environments (Virtualization)

- **Virtualization** vs. **Emulation** (both are SW)
 - **OS to OS** vs. **CPU to CPU**
- Use cases involve laptops and desktops running multiple OSes for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSes without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers
- VMM can run natively, in which case they are also the host
 - There is no general purpose host then (VMware, ESX and Citrix XenServer)

Computing environments (Virtualization)



(a)

achine


Computing environments (Cloud Computing)

- Delivers **computing**, **storage**, even **apps** as a **service** across a network
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage
- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more **applications** available via the Internet (i.e., word processor)
 - Platform as a Service (**PaaS**) – **software stack** ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – **servers** or **storage** available over Internet (i.e., storage available for backup use)

Computing environments (Cloud Computing)

➤ Cloud computing environments composed of **traditional OSe**, plus **VMMs**, plus **cloud management** tools

- Internet connectivity **requires** security like **firewalls**
- Load balancers **spread traffic** across **multiple** applications



Computing environments (Real-Time Embedded Systems)

- Real-time embedded systems most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - Use expanding
- Many other special computing environments as well
 - Some have OSes, some perform tasks without an OS
- Real-time OS has well-defined fixed time constraints
 - Processing **must** be done within **constraint** (soft, firm, hard)
 - Correct operation only if constraints met



Open-source operating systems

- Operating systems made available in **source-code** format rather than just binary **closed-source**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - **<http://www.virtualbox.com>**)
 - Use to run guest operating systems for exploration
- All book materials are available at: **www.os-book.com**

Questions?

