



Amirkabir University of Technology
(Tehran Polytechnic)
Department of Computer Engineering

Process Synchronization

(همگام سازی فرایندها)

Hamid R. Zarandi
h_zarandi@aut.ac.ir

Motivation

- **Cooperating** process/thread:
 - the one that can affect or be affected by other processes executing in system.
 - **Processes, threads**
- Processes can execute **concurrently**
 - May be interrupted at any time, **partially completing execution**
- Problem: **Data inconsistency** (ناسازگاری داده)
 - It may occur in the case of concurrent access to shared data
- **How to solve?**
 - **Orderly execution of cooperating processes that share a logical address space**

One example!

- A solution to **consumer-producer** problem that fills *all* the buffers.
 - We can have an integer **counter** that keeps track of the **number of full buffers**.
 - Initially, **counter** is **set to 0**.
 - It is **incremented** by the **producer** after it produces a new buffer
 - It is **decremented** by the **consumer** after it consumes a buffer.

Circular buffer & producer-consumer problem

```

#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;

```

Producer

```

item next_produced;
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter ++;
}

```

Consumer

```

item next_consumed;
while (true) {
    while (counter == 0)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter --;

    /* consume the item in next consumed */
}

```

Race condition



➤ **counter++** could be implemented as

<code>register1 = counter</code>	<code>MOV AX, [100]</code>
<code>register1 = register1 + 1</code>	<code>ADD AX, 1</code>
<code>counter = register1</code>	<code>MOV [100], AX</code>

➤ **counter--** could be implemented as

<code>register2 = counter</code>	<code>MOV BX, [100]</code>
<code>register2 = register2 - 1</code>	<code>ADD BX, 1</code>
<code>counter = register2</code>	<code>MOV [100], BX</code>

➤ Consider this execution interleaving with “count = 5” initially:

S0: producer execute <code>register1 = counter</code>	{register1 = 5}
S1: producer execute <code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute <code>register2 = counter</code>	{register2 = 5}
S3: consumer execute <code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute <code>counter = register1</code>	{counter = 6}
S5: consumer execute <code>counter = register2</code>	{counter = 4}

Another Race condition

➤ Invoking *echo()* procedure:

```
void echo()  
{  
    chin = getchar();  
    chout = chin;  
    putchar(chout);  
}
```

Process P1

-
- chin = getchar();
-
- chout = chin;
- putchar(chout);
-
-

Process P2

-
-
- chin = getchar();
- chout = chin;
-
- putchar(chout);
-

➤ Same problem exists on:

- Multiprogramming environment
- Multiprocessing environment
- Distributed processing environment

Other examples?

Have you ever seen other examples?

Definition

➤ Race condition

- Several process access and manipulate the same data **concurrently**
- Outcomes of the execution **depends** on the **order** in which the access take place

➤ How to remove **Race Condition**?

- **Serial execution**

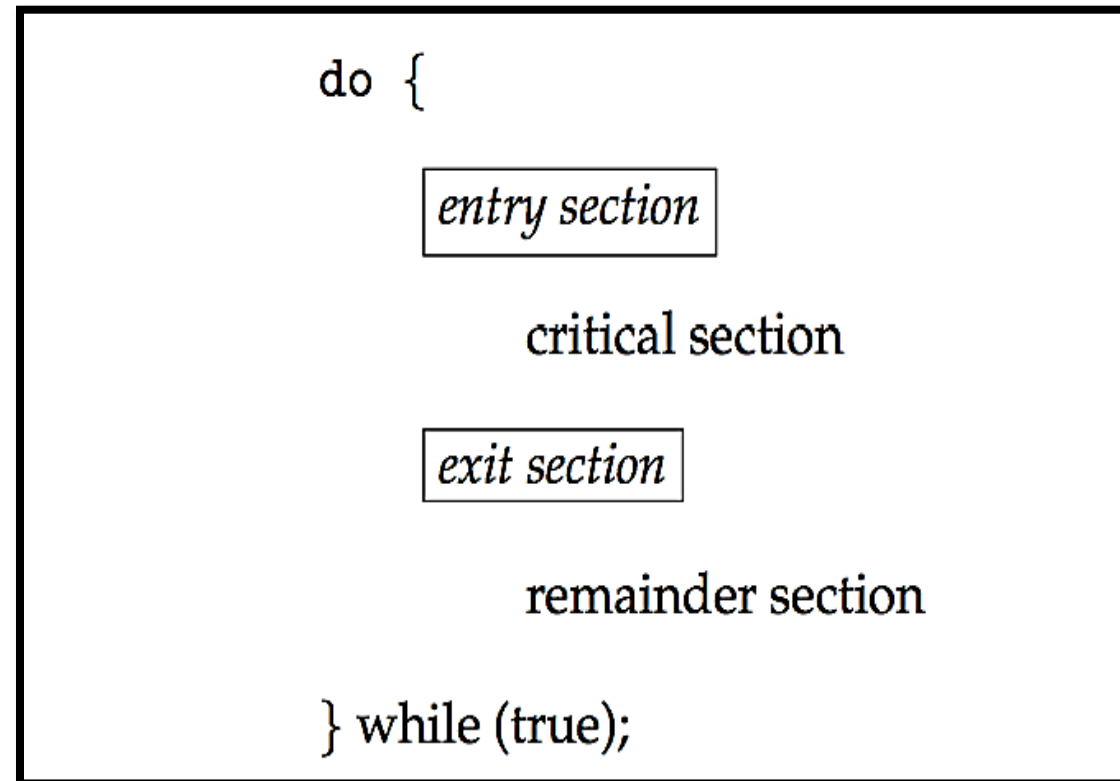
Critical Section Problem

Critical section problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process **must ask permission** to enter critical section in **entry section**, may follow critical section with **exit section**, then remainder section

Critical section

➤ General structure of process P_i



Requirements to solutions

➤ Mutual exclusion (انحصار متقابل)

- If process P_i is executing in its critical section, then **no other processes** can be executing in their critical sections

➤ Progress (پیشرفت)

- If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the **selection of the processes** that will enter the critical section next **cannot be postponed indefinitely**

➤ Bounded waiting (انتظار محدود)

- A **bound** must exist on the **number of times that other processes are allowed** to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes

Preemption definition

➤ Preemption (قبضه ای - قبضه شدنی)

- The act of temporarily interrupting a [task](#) being carried out by a [computer system](#), without requiring its cooperation, and with the intention of resuming the task at a later time [[wiki](#)]

Handling critical-section by OS

➤ Two approaches, depend on type of OS kernels

○ Preemptive

- Allows preemption of process when running in kernel mode
- Difficult to design in SMP architectures (why?)

○ Non-preemptive

- Runs until exits kernel mode, blocks, or voluntarily yields CPU
 - ✓ Essentially free of race conditions in kernel mode (why?)

➤ Which one

○ is responsive?

○ is suitable for real-time programming?

1) Peterson's solution

- A classic **SW** solution
- **No guarantees** in correct working of the method
 - Correctness depends on computer architecture
 - **Atomic** instructions are needed (**which & where?**)
- **Good algorithm!**
- **Shared variables**
 - `int turn; /* whose turn is */`
 - `Boolean flag[2] /* who enters the critical-section */`

Peterson algorithm for P_i

$(P_i, P_j) = (P_0, P_1)$

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = false;  
        remainder section  
} while (true);
```

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)

2) Hardware solution

- Some hardwares support implementing the critical section code!
- All solutions are based on idea of **locking**
 - Protecting critical regions via locks
- **Uniprocessors** – **could disable interrupts**
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - Operating systems using this not broadly scalable
- **Multiprocessors** – **provide special atomic hardware instructions**
 - **Atomic** = non-interruptible
 - Either
 - test memory word and set value
 - swap contents of two memory words

Hardware solution for critical section

```
do {  
    acquire lock  
        critical section  
    release lock  
        remainder section  
} while (TRUE);
```

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)

test_and_set instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;    /* old value */
}
```

1. Executed **atomically**
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to "TRUE".

Hardware solution using *test_and_set()*

- Shared Boolean variable lock, initialized to FALSE

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */  
} while (true);
```

compare_and_swap instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return temp;    /* old value */  
}
```

1. Executed **atomically**
2. Returns the original value of passed parameter "value"
3. Set the variable "value" the value of the passed parameter "new_value" but only if "value" == "expected". That is, the swap takes place only under this condition.

Hardware solution using *compare_and_swap()*

➤ Shared integer “lock” initialized to 0;

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
} while (true);
```

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)

Bounded-waiting mutual exclusion with test_and_set

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;

    /* critical section */

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;

    /* remainder section */

} while (true);
```

3) OS solution!: Mutex locks

- Previous solutions are **complicated** and generally **inaccessible** to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is **mutex** lock (*mutual exclusions*)
- Protect a critical section by first **acquire()** a lock then **release()** the lock
 - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic**
 - Usually implemented via **hardware atomic instructions**
- But this solution requires **busy waiting**
 - This lock therefore called a **spinlock**

acquire() and release()

```
➤ acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

```
➤ release() {  
    available = true;  
}
```

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)

What is the main problem of all mentioned methods?

Busy waiting!

4) Semaphore

- **Synchronization tool** that provides more **sophisticated ways** (than **Mutex** locks) for process to synchronize their activities
- Semaphore S – **integer variable**
- Can only be accessed via two indivisible (**atomic**) operations

wait() and **signal()**

```
wait(S)
{
    while (S <= 0)
        ; // busy wait
    S--;
}
```

```
signal(S)
{
    S++;
}
```

No busy waiting in Semaphore

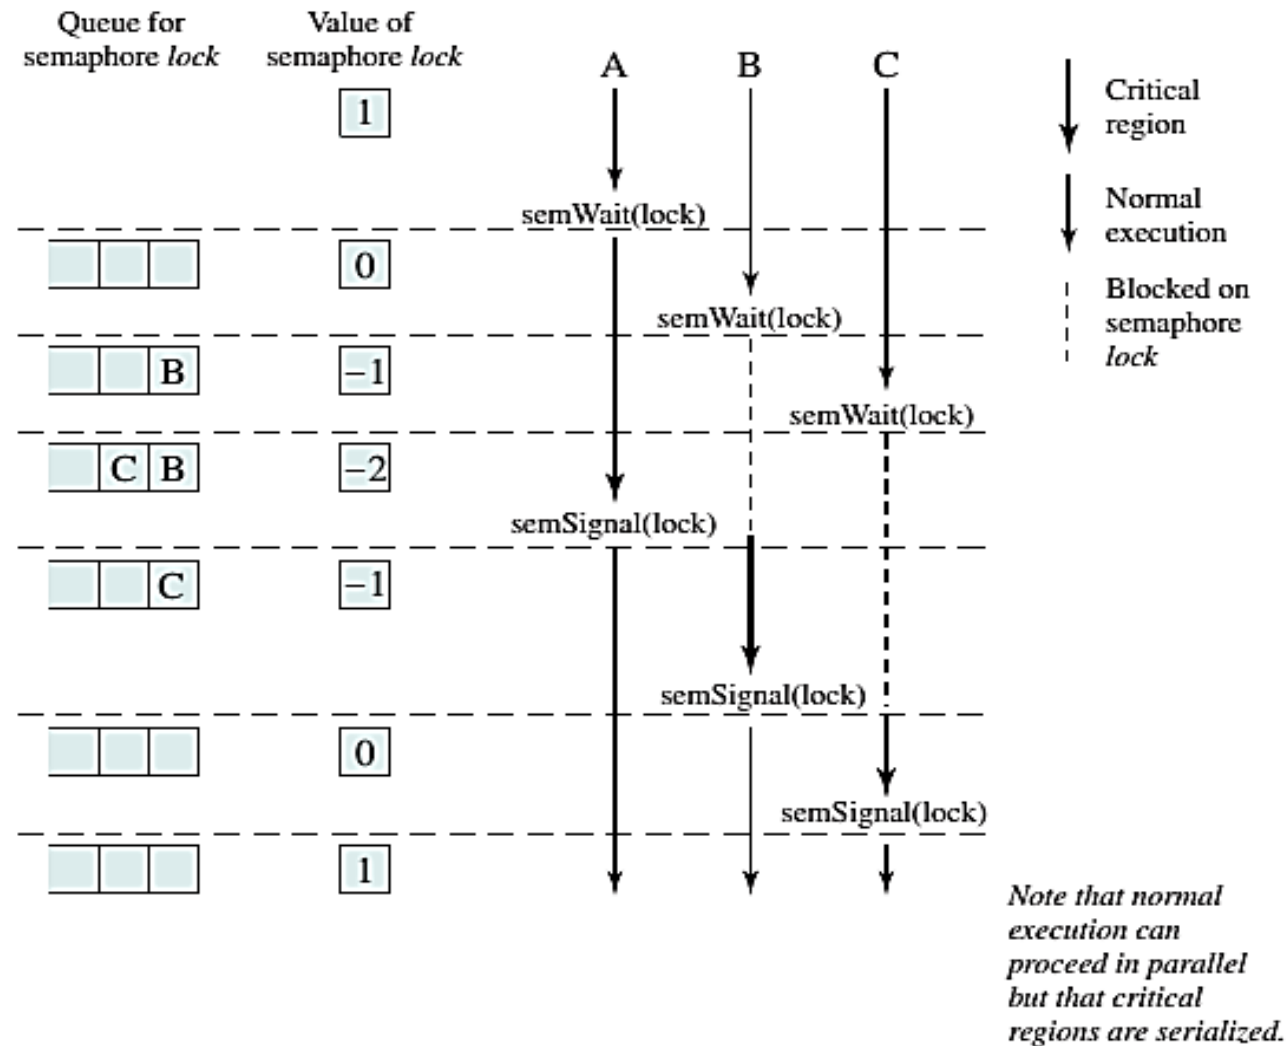
- Have a FIFO queue for waiting process

```
typedef struct{
    int value;
    struct process *list;
} semaphore;
```

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}
```

```
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

Accessing shared data by Semaphore



Types of semaphore

➤ Types

- Binary semaphore (same as mutex lock)
- Counting semaphore (suitable for managing number of resources)

➤ Can solve various synchronization problems

➤ Example:

- Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to **zero**

P1:

```
 $S_1$  ;  
signal ( synch ) ;
```

P2:

```
wait ( synch );  
 $S_2$  ;
```

Semaphore points

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time (**why?**)
 - `wait()` and `signal()` must be **atomic!**
 - `wait()` and `signal()` generate a **Critical Section Problem!**
 - **How to solve?**
 - Uniprocessors
 - ✓ **Disabling interrupts**
 - SMP (Multiprocessors)
 - ✓ **Disabling interrupts** (**bad performance effect**)
 - ✓ Other methods: **`compare_and_swap()`** and **`spinlock`** (**is it good to have busy waiting?**)

Two implementations of semaphores

```
semWait(s)
{
    while (compare_and_swap(s.flag, 0 , 1) == 1)
        /* do nothing */;
    s.count--;
    if (s.count < 0) {
        /* place this process in s.queue*/;
        /* block this process (must also set
s.flag to 0) */;
    }
    s.flag = 0;
}

semSignal(s)
{
    while (compare_and_swap(s.flag, 0 , 1) == 1)
        /* do nothing */;
    s.count++;
    if (s.count <= 0) {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
    s.flag = 0;
}
```

(a) Compare and Swap Instruction

```
semWait(s)
{
    inhibit interrupts;
    s.count--;
    if (s.count < 0) {
        /* place this process in s.queue */;
        /* block this process and allow inter-
rupts*/;
    }
    else
        allow interrupts;
}

semSignal(s)
{
    inhibit interrupts;
    s.count++;
    if (s.count <= 0) {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
    allow interrupts;
}
```

(b) Interrupts

Problems with semaphores

➤ Be careful in the usage

- Deadlock, Starvation, Priority inversion

➤ Starvation

- LIFO queue

P_0

```
wait(S);  
wait(Q);  
...  
signal(S);  
signal(Q);
```

P_1

```
wait(Q);  
wait(S);  
...  
signal(Q);  
signal(S);
```

➤ Priority Inversion – Scheduling problem when lower-priority process holds a lock needed by higher-priority process

- Example: $L(R) < M < H(R)$

➤ Solved via priority-inheritance protocol

Classic synchronization problems

- The bounded-buffer problem
- The readers-writers problem
- The dining-philosophers problem

How can semaphore solve these problems?

The bounded-buffer problem

```
int n;  
semaphore mutex = 1;  
semaphore empty = n;  
semaphore full = 0;
```

Producer

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
  
    wait(empty);  
    wait(mutex);  
  
    ...  
    /* add next produced to the buffer */  
    ...  
  
    signal(mutex);  
    signal(full);  
  
} while (true);
```

Consumer

```
do {  
    wait(full);  
    wait(mutex);  
    ...  
  
    /* remove an item from buffer to next_consumed */  
    ...  
  
    signal(mutex);  
    signal(empty);  
  
    ...  
  
    /* consume the item in next consumed */  
    ...  
  
} while (true);
```

The readers-writers problem

```
semaphore rw_mutex = 1;
semaphore mutex = 1;
int read_count = 0;
```

Readers

Writers

```
do {
    wait(rw_mutex);

    ...

    /* writing is performed */

    ...

    signal(rw_mutex);

} while (true);
```

```
do {

    wait(mutex);
    read_count++;
    if (read_count == 1)
        wait(rw_mutex);
    signal(mutex);

    ...

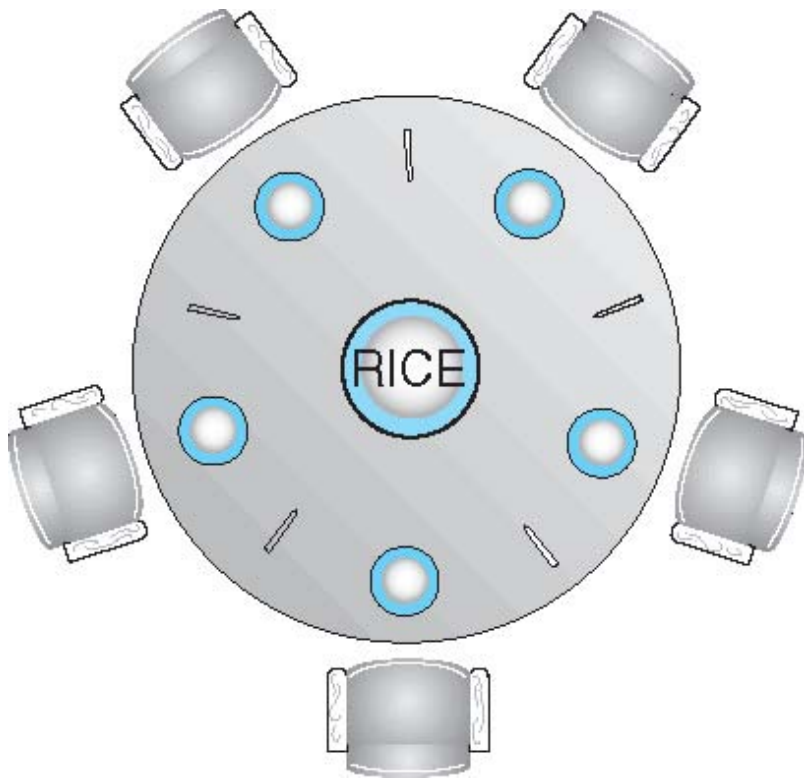
    /* reading is performed */

    ...

    wait(mutex);
    read count--;
    if (read_count == 0)
        signal(rw_mutex);
    signal(mutex);
} while (true);
```

The dining-philosophers problem

➤ Thinking and eating alternatively



```
semaphore chopstick[5];
```

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5] );  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
} while (TRUE);
```

Any problem?

Other problems with semaphore

➤ Problems with **bad** usage

```
signal(mutex);  
...  
critical section  
...  
wait(mutex);
```

```
wait(mutex);  
...  
critical section  
...  
wait(mutex);
```

```
...  
critical section  
...  
wait(mutex);
```

```
wait(mutex);  
...  
critical section  
...
```

➤ **Deadlock** and **starvation** are possible.

5) Monitor

- A **high-level abstraction** that provides a convenient and effective mechanism for process synchronization
- Only **one process** may be active within the monitor at a time

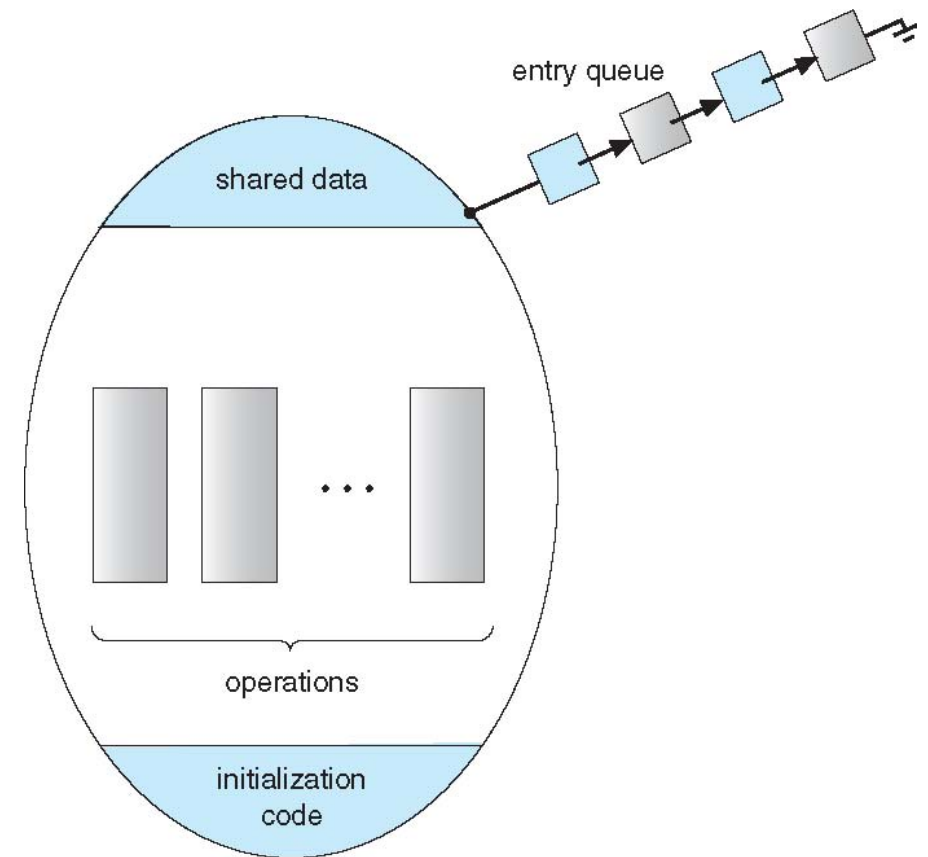
```

monitor monitor-name
{
    // shared variable declarations

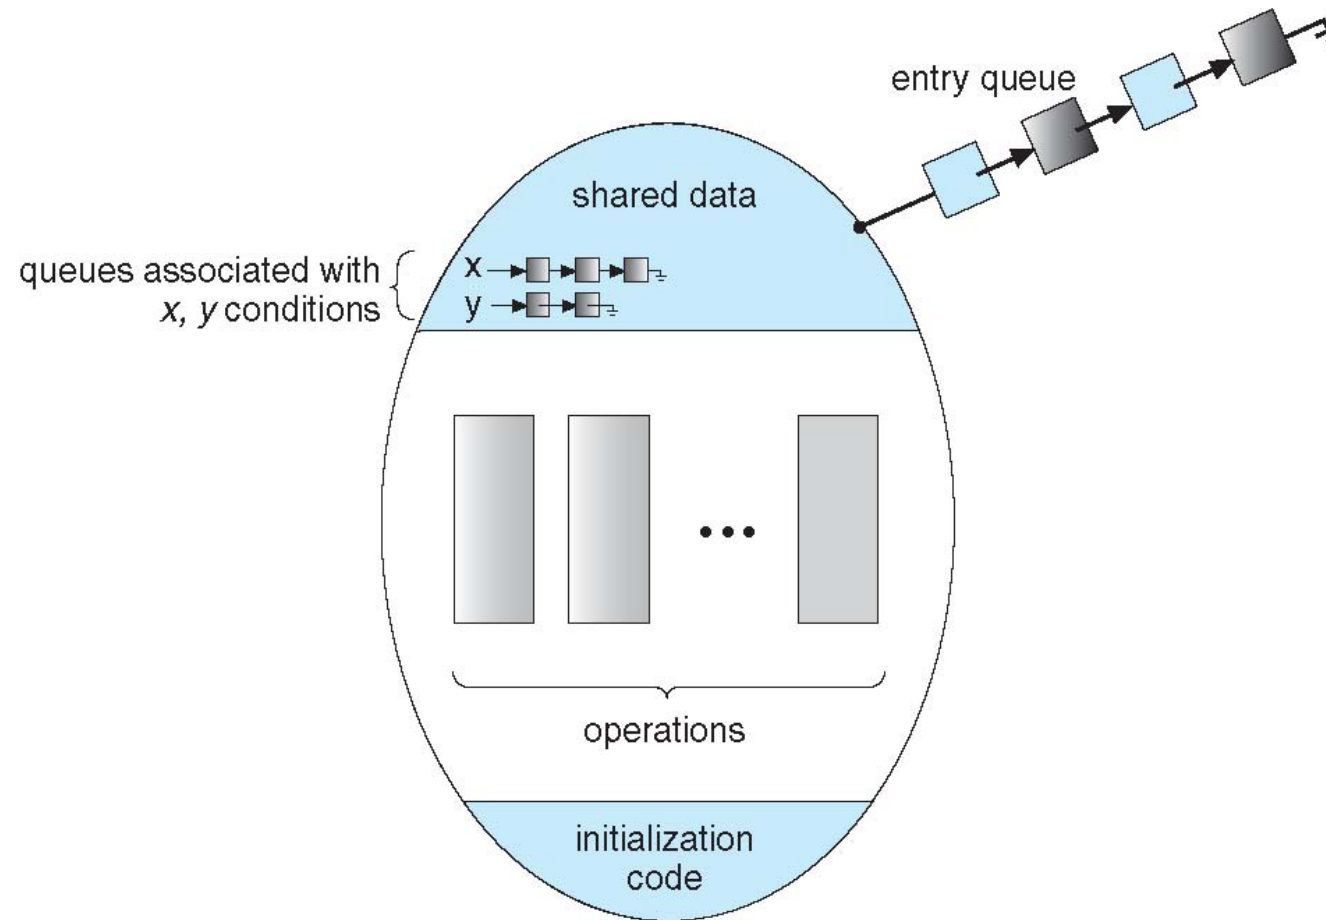
    procedure P1 (...) { ... }

    procedure Pn (...) {.....}

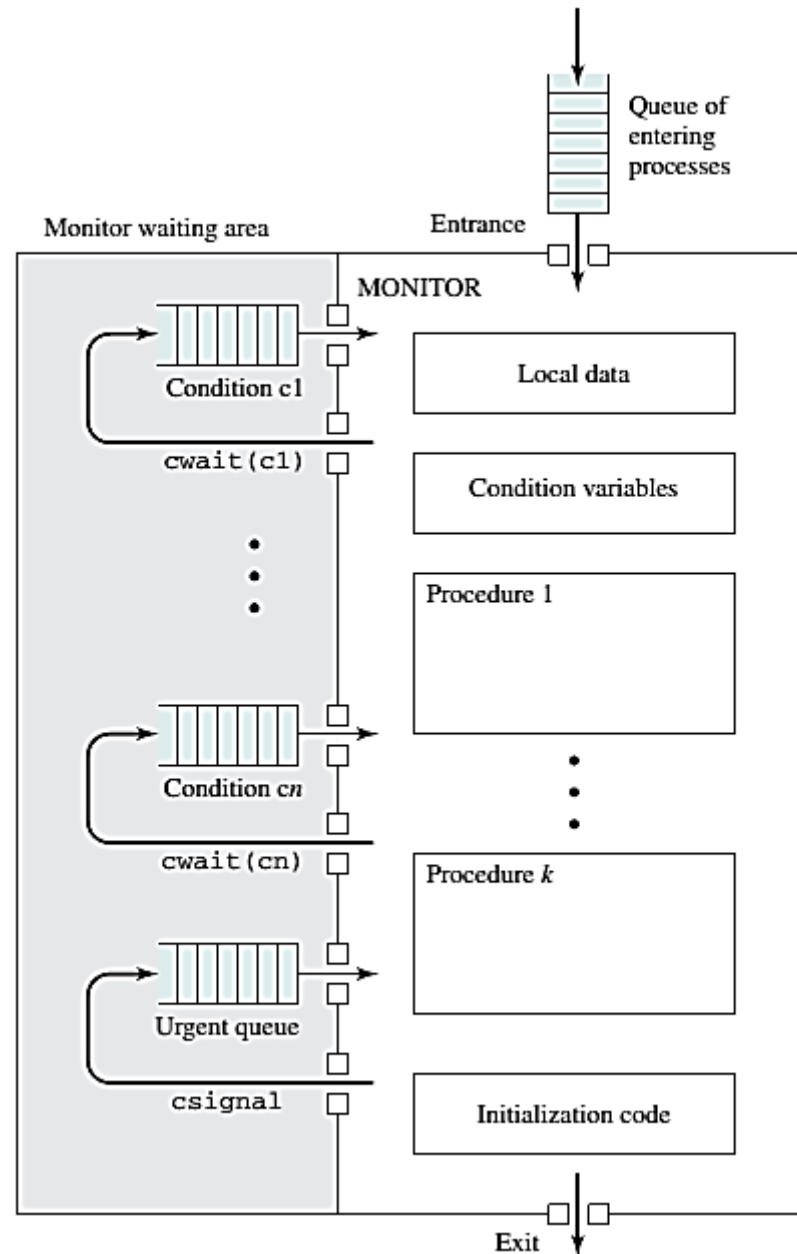
    Initialization_Code (...) { ... }
}
  
```



Monitor (with *condition variables*)



Structure of a Monitor



The dining-philosophers problem

```
monitor DiningPhilosophers
```

```
{
```

```
    enum { THINKING; HUNGRY, EATING) state [5] ;
```

```
    condition self [5];
```

```
    void pickup (int i) {
```

```
        state[i] = HUNGRY;
```

```
        test(i);
```

```
        if (state[i] != EATING)
```

```
            self[i].wait();
```

```
    }
```

```
    void putdown (int i) {
```

```
        state[i] = THINKING;
```

```
        // test left and right neighbors
```

```
        test((i + 4) % 5);
```

```
        test((i + 1) % 5);
```

```
    }
```

```
    void test (int i) {
```

```
        if ((state[(i + 4) % 5] != EATING) &&
```

```
            (state[i] == HUNGRY) &&
```

```
            (state[(i + 1) % 5] != EATING) ) {
```

```
                state[i] = EATING;
```

```
                self[i].signal();
```

```
        }
```

```
    }
```

```
    initialization_code() {
```

```
        for (int i = 0; i < 5; i++)
```

```
            state[i] = THINKING;
```

```
    }
```

```
}
```

The dining-philosophers problem



```
DiningPhilosophers.pickup(i);
```

EAT

```
DiningPhilosophers.putdown(i);
```

Any problem?

No deadlock

Starvation is possible

Solving bounded-buffer using a Monitor

```

/* program producerconsumer */
monitor boundedbuffer;
char buffer [N];                                /* space for N items */
int nextin, nextout;                             /* buffer pointers */
int count;                                       /* number of items in buffer */
cond notfull, notempty;                        /* condition variables for synchronization */
void append (char x)

{
    if (count == N) cwait(notfull);             /* buffer is full; avoid overflow */
    buffer[nextin] = x;
    nextin = (nextin + 1) % N;
    count++;
    /* one more item in buffer */
    csignal (notempty);                        /*resume any waiting consumer */
}

void take (char x)
{
    if (count == 0) cwait(notempty);            /* buffer is empty; avoid underflow */
    x = buffer[nextout];
    nextout = (nextout + 1) % N;
    count--;
    /* one fewer item in buffer */
    csignal (notfull);                        /* resume any waiting producer */
}

/* monitor body */
/* buffer initially empty */
nextin = 0; nextout = 0; count = 0;

```

```

void producer()
{
    char x;
    while (true) {
        produce(x);
        append(x);
    }
}

void consumer()
{
    char x;
    while (true) {
        take(x);
        consume(x);
    }
}

void main()
{
    parbegin (producer, consumer);
}

```

Points to monitor

- Monitors can be implemented by semaphores (See the textbook).
- OSes support
 - Monitor, semaphore, spinlock, mutex
 - Examples
 - Solaris
 - Windows
 - Linux
 - Pthreads
- Alternative approaches
 - Transactional Memory
 - OpenMP
 - Functional Programming Languages

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

Questions?

